

Introduction

Dr. RAJIB MALL

Professor

**Department Of Computer Science &
Engineering**


IIT Kharagpur.

Course Plan

- Introduction
- Basics of task scheduling
 - Cyclic executives
 - RMA and EDF
- Resource Sharing
- Scheduling on Multiprocessors
- POSIX-RT
- Commercial RTOS

Text Book

- R. Mall, Real-Time Systems, Pearson, 2008.
- To be supplemented with other materials



Reference Books


- Jane Liu, Real-Time Systems, Pearson, 2000
- C. Krishna and K. Shin, Real-Time Systems, McGraw-Hill, 2000

What is Real-Time?

- Real-time is a quantitative notion of time measured using a physical clock.
 - **Example:** After a certain event occurs (temperature exceeds 500 degrees) the corresponding action (coolant shower) must complete within 100mSec.
- This is in contrast to the qualitative notion of time:
 - Expressed using notions such as **before, after, sometime, eventually**, etc.

Real-Time Systems

- Characterized by time-constrained response to events.
- Often are:
 - Embedded
 - Safety-critical
- RTOS helps applications to meet their deadlines:
 - Task scheduling is the primary mechanism for making applications meet their respective deadlines.




Why Surge in Embedded Applications?

- Trend of reducing cost of computers:
 - Processors
 - Memory
- Flexibility due to Internet
- Reducing power consumption
- Reducing size
- Increasing Processing power and hardware and software reliability

Important Characteristics

- An embedded system responds to events.



Example: An Automobile airbag system.

When the airbag's motion sensors detect a collision, the system needs to respond by deploying the airbag within 10ms or less.

- or the system fails!

How are RTOS Different from other OS?

- An embedded system responds to external inputs:
 - If response is late, the system fails.
- General purpose OS:
 - Minimizes response time and ensures fairness
- Real-time OS:
 - Helps tasks meet their deadline.

Characteristics of A Real-Time System

- **Real-time:**

- At least tasks have real-time constraints, e.g a Deadline.

- **Correctness Criterion:**

- Results should be logically correct,
- **And also within the stipulated time.**

Why Have an OS in an Embedded Device?

- Support for:
 - Multitasking, scheduling, and synchronization
 - Timing aspects.
 - Memory management
 - File systems
 - Networking
 - Graphics displays
 - Interfacing wide range of I/O devices
 - Scheduling and buffering of I/O operations
 - Security and power Management

Embedded OS

- **Example:** A smartphone operating system contains over five million lines of code!
- Projects will hardly have the time and funding:
 - To develop all of this code on their own!
- Typical Embedded OS license fees are less than even Rs100 per device --- lower than a desktop OS
- Some very simple low-end devices might not need an OS:
 - But devices are getting more complex.

Types of Real-Time Systems


- Real-time systems are different from traditional systems:
 - Tasks have deadlines associated with them.
- Classified largely based on the consequence of not meeting deadline:
 - Hard real-time systems
 - Soft real-time systems
 - Firm real-time systems

Hard Real-Time Systems

- If a deadline is not met:
 - The system is said to have failed.
- The task deadlines are of the order of micro or milliseconds.
- Many hard real-time systems are safety-critical.
- Examples:
 - Industrial control applications
 - On-board computers
 - Robots

Firm Real-Time Systems

- If a deadline is missed occasionally, the system does not fail:
 - The results produced by a task after the deadline are ignored.




Firm Real-Time Systems

- Examples:

- A video conferencing application
- A telemetry application
- Satellite-based surveillance applications


Soft Real-Time Systems

- If a deadline is missed, the system does not fail:
 - The utility of a result decreases with time after the deadline.
 - If several tasks miss deadline, then the performance of the system is said to have degraded.



Soft Real-Time Systems

- Use probabilistic requirements on deadline.
- For example, 99% of time deadlines will be met.



Soft Real-Time Systems

- Examples:
 - Railway reservation system
 - Web browsing
 - In fact, all interactive applications

Types of Tasks

- **Periodic:**
 - Recur according to a timer
 - A vast majority all real-time tasks are periodic
- **Aperiodic:**
 - Recur randomly and are soft real-time tasks
- **Sporadic:**
 - Recur randomly, but hard real-time tasks

Timing Constraints

- A timing constraint:
 - Defined with respect to some event.
- An event:
 - Occurs at an instant of time
 - Generated either by the system or its environment

Real-Time Tasks


- Real-time tasks get generated due to certain event occurrences:
 - Either internal or external events.
- Example:
 - A task may get generated due to a temperature sensor sensing high-level.
- When a task gets generated:
 - It is said to be released or arrived .

Real-Time Task Scheduling

- Essentially refers to the order in which the various tasks are to be executed.
- It is the primary means adopted by an operating system to meet task deadlines.
- Obviously, scheduler is a very important component of every RTOS.

Real-Time Workload

- **Job:**
 - A unit of work
 - A computation, a file read, a message transmission, etc
 - **A task instance**
- **Task:** a sequence of similar jobs




Task Instance (Job)

- A task typically recurs a large number of times:
 - Each time triggered by an event
 - Each time a task recurs, an instance of the task is said to have been generated or released.
- The i th time a task T recurs:
 - Job or Task instance T_i is said to have arrived.


Relative and Absolute Deadlines

- **Absolute deadline:**
 - Counted from time 0.
- **Relative deadline:**
 - Counted from time of occurrence of task.



Response Time

- Duration between task release time and task completion time.
- Release time
 - The time of occurrence of the event generating the task.
- Completion time
 - Results produced by the task

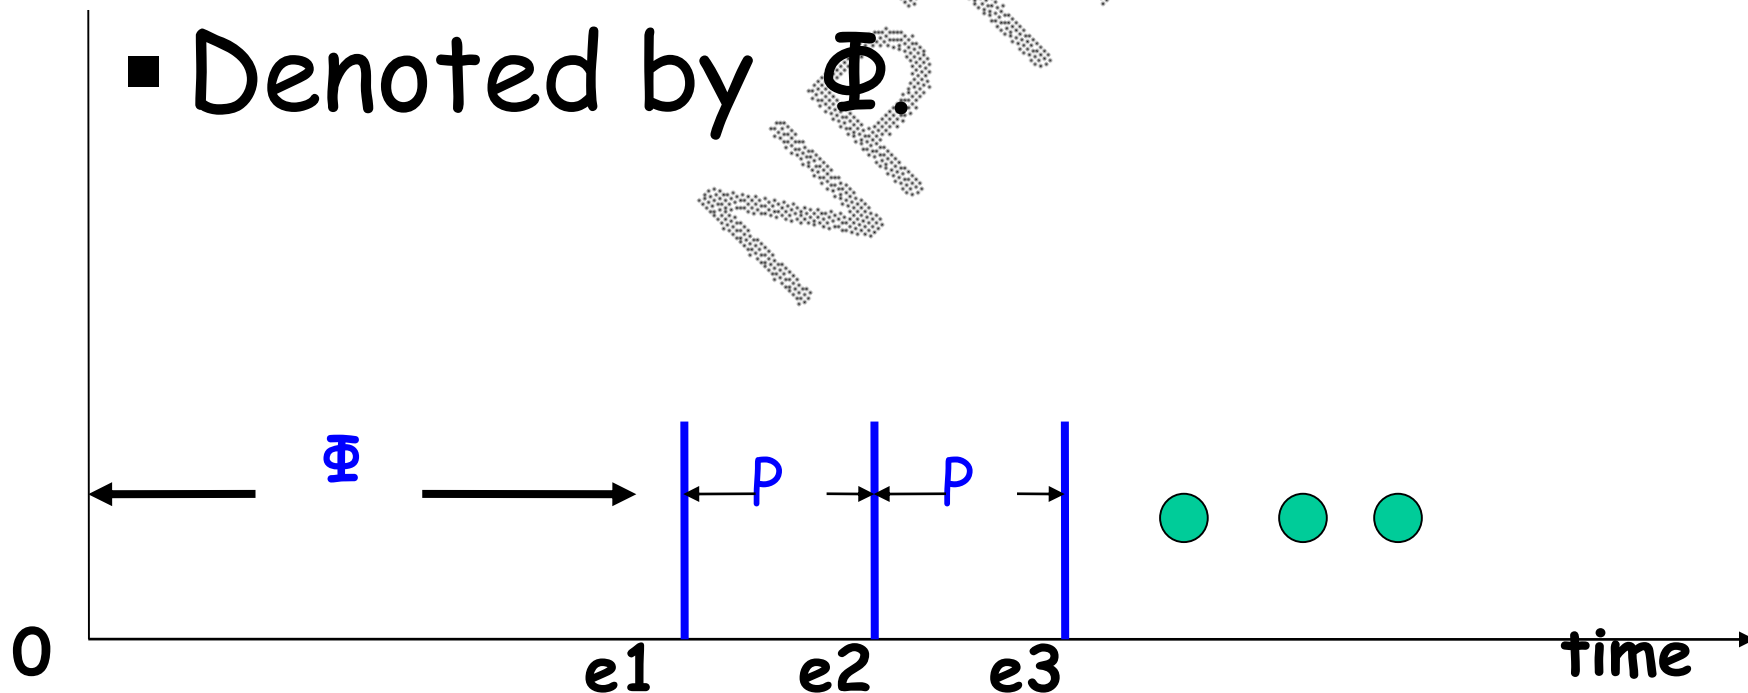


Response Time


- For soft real-time tasks:
 - The response time needs to be minimized.
- For hard real-time tasks:
 - As long as the task completes within its deadline,
 - No advantage of completing it any early.

Phase of a Periodic Task

- Phase for a periodic task:
 - The time from 0 till the occurrence of the first instance of the task.
 - Denoted by Φ .



Phase Example



- The track correction task starts 2000 mSecs after the launch of the rocket:
 - Periodically recurs every 50 milli Seconds then on.
 - Each instance of the task requires a processing time of 8 mSecs and its relative deadline is 50 mSecs.

Few Task Scheduling Terminologies

- Valid Schedule:

- At most one task is assigned to a processor at a time.
- No task is scheduled before it is ready.
- Precedence and resource constraints of all tasks are satisfied.

- Feasible Schedule:

- Valid schedule is one in which all tasks meet their respective time constraints₄₂

Scheduling Terminologies

▪ Proficient Scheduler:

- A scheduler $S1$ is more proficient compared to another Scheduler $S2$:
 - If whichever tasks that $S2$ can feasibly schedule so can $S1$, but not vice versa.

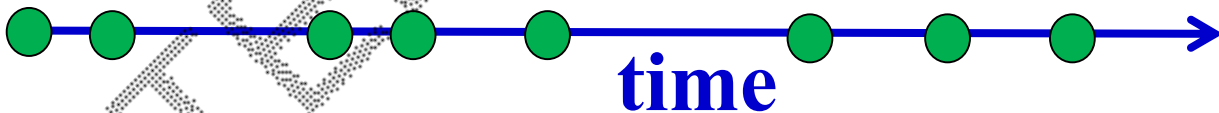
▪ Equally proficient schedulers:

- If a task set scheduled by one can also be scheduled by the other and vice versa.

Scheduling Terminologies

- Optimal Scheduler:
 - An optimal scheduler can feasibly schedule any task set that can be scheduled by any other scheduler.

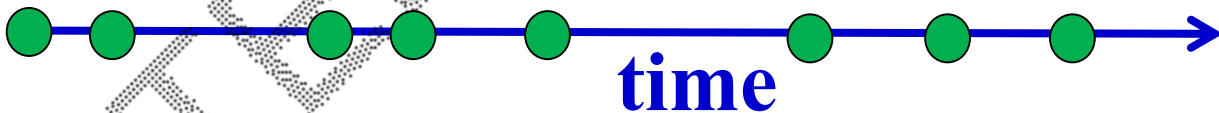
Scheduling Points

- **At these points on time line:**
 - Scheduler makes decision regarding which task to be run next.
- **Clock-driven:** 
 - Scheduling points are defined by interrupts from a periodic timer.
- **Event-driven:**
 - Scheduling points defined by task completion and generation events.

Real-Time Task Scheduling

- Significant amount of research has been carried out to develop schedulers for real-time tasks:
 - Schedulers for uniprocessors
 - Schedulers for multiprocessors and distributed systems.

Scheduling Points

- **At these points on time line:**
 - Scheduler makes decision regarding which task to be run next.
- **Clock-driven:** 
 - Scheduling points are defined by interrupts from a periodic timer.
- **Event-driven:**
 - Scheduling points defined by task completion and generation events.

Real-Time Task Scheduling

- Significant amount of research has been carried out to develop schedulers for real-time tasks:
 - Schedulers for uniprocessors
 - Schedulers for multiprocessors and distributed systems.

Task Scheduling on Uniprocessors

- Focus of much research during the 1970s and 80s.
- Real-time task schedulers can be broadly classified into:
 - Clock-driven
 - Event-driven

Summary of Schedulers

- Endless Loop

- No Tasks, Polled

- Simple
Cyclic Executive

- Single frequency

- Multi-rate Cyclic
Executive

- Multiple frequencies

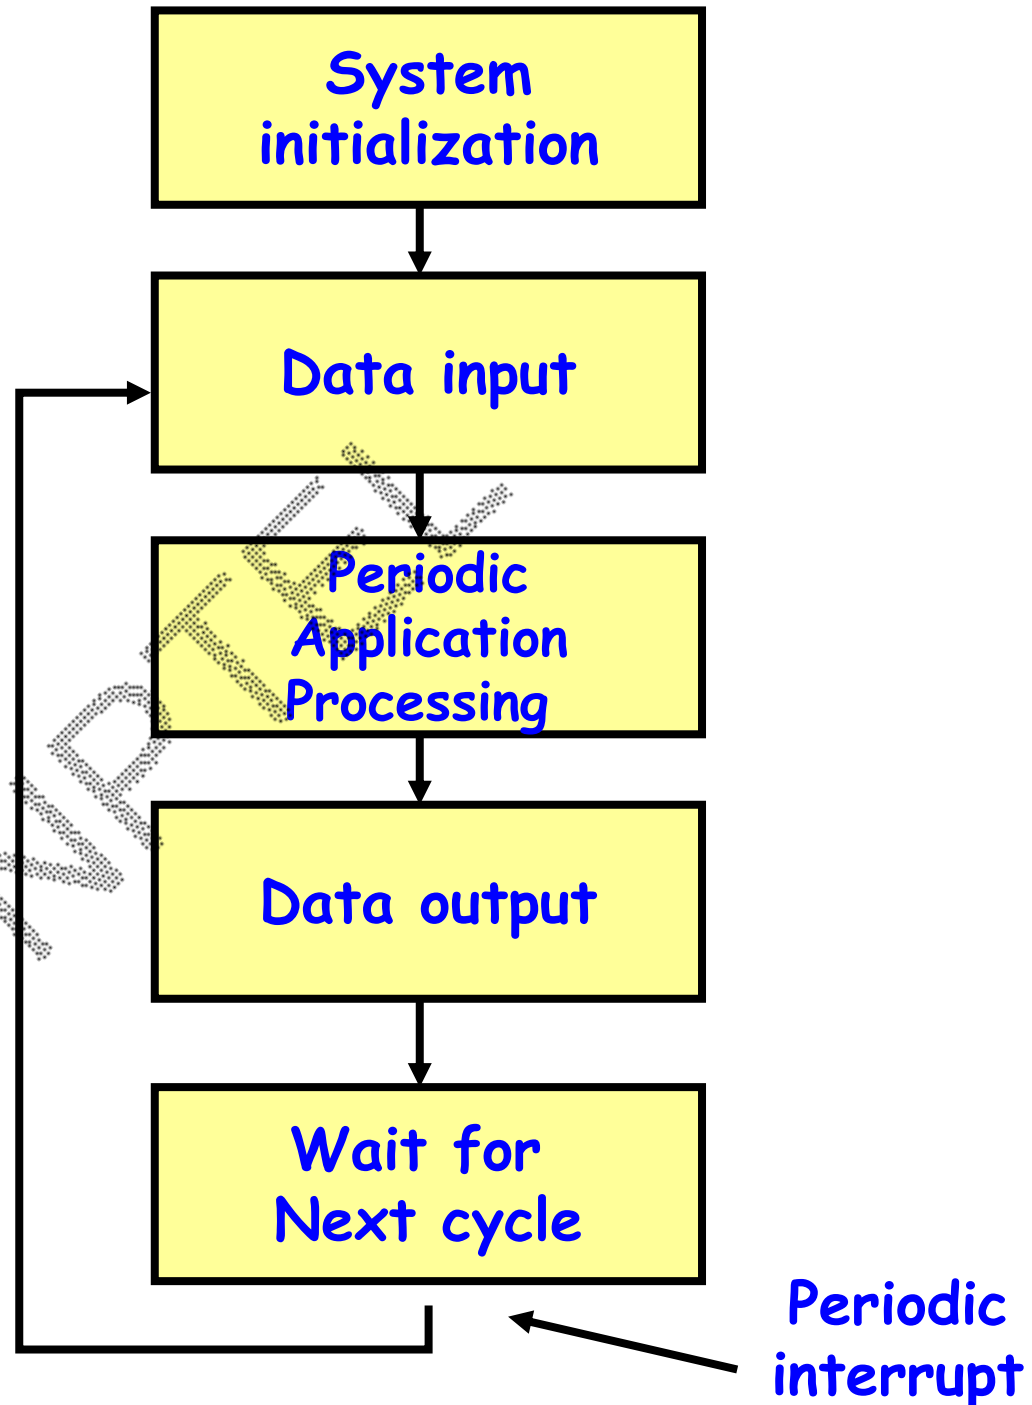
- Priority-based
Preemptive Scheduler

- Interrupt driven

-

Cyclic Executives

Simple Cyclic Executive




Cyclic Executive

- No actual processes exist:
 - Each cycle only involves a sequence of procedure calls.
- The procedures share a common address space and share data.
 - This data need not be protected (via a semaphore, for example) because “processes” are not preempted.
- All “process” periods are multiples of the cycle time.

Clock-Driven Scheduling: Basics

- Decision regarding which job to run next is made only at clock interrupt instants:



- Interval timers determine the scheduling points.
- Which task to be run when and for how long is stored in a table.

Clock-Driven Scheduling

- Round robin scheduling:
- Popularly used:
 - Basic Timer-Driven Scheduler (Table-driven)
 - Cyclic Scheduler

Clock-Driven Schedulers

- Also called:
 - Offline schedulers
 - Static schedulers
- Used extensively in embedded applications:
 - Table driven schedulers
 - Cyclic schedulers

Clock-Driven Scheduling

- Used in low cost applications:

- **Pro:**

- **Compact:** Require very little storage space
- **Efficient:** Incur very little runtime overhead.

- **Con:**

- **Inflexible:** Very difficult to accommodate aperiodic or sporadic tasks.

- The simplest is table-driven scheduler.

Basic Table-Driven Scheduler

```
const int SchedTableSize= 10;
```

```
timer_handler () {
```

```
    int next_time; task current;
```

```
    current = SchedTable[entry].tsk;
```

```
    entry = (entry+1) % SchedTableSize;
```

```
    next_time = Table[entry].time + gettime();
```

```
    set_timer(next_time);
```

```
    execute_task(current);
```

```
    return;
```

```
}
```

(ScheduleTable)

T_1	$t(T_1)$
..	...
T_n	$t(T_n)$

Schedule Table

Task	Time
T1	90
T2	120
T3	75
T4	225
T5	50

Disadvantage of Table-Driven Schedulers

- When the number of tasks are large:
 - Requires setting the timer large number of times.
 - The overhead is significant:
 - Remember that a task instance runs only for a few milli or microseconds.

Cyclic Schedulers

- Cyclic schedulers are very popular:
 - Extensively being used.
- Many tiny embedded applications have severe constraints on memory and processing power:
 - Cannot even host a microkernel RTOS
 - Use cyclic schedulers.

Cyclic Schedulers

- For scheduling n periodic tasks:
 - The schedule is stored in a table.
 - Repeated forever.
 - The designer needs to develop a schedule for what period?
 - $LCM(P_1, P_2, \dots, P_n)$

Cyclic Schedulers


- Cyclic schedulers are very popular:
 - Extensively being used.
- Many tiny embedded applications have severe constraints on memory and processing power:
 - Cannot even host a microkernel RTOS
 - Use cyclic schedulers.

Cyclic Schedulers

- For scheduling n periodic tasks:
 - The schedule is stored in a table.
 - Repeated forever.
 - The designer needs to develop a schedule for what period?
 - $LCM(P_1, P_2, \dots, P_n)$

Cyclic Schedulers


- The schedule is Precomputed and stored for one **major cycle**:
 - This schedule is repeated.
- A major cycle is divided into:
 - One or more **minor cycles** (frames).
- Scheduling points for a cyclic scheduler:



- Occur at the beginning of frames.

Cyclic Scheduler Basics

- Scheduling decisions are only made at frame boundaries.




- Exact start and completion time of jobs within frame is not known
- But must know the max computation time
- Jobs are allocated to specific frames
- Major cycle is also called a **Hyperperiod**.

Cyclic Scheduler Basics

- If a schedule can not be found for the set of predefined jobs:
 - Then these are divided into **job slices**.
 - Essentially, divide a job into a sequence of smaller jobs.

Major Cycle

- In each major cycle:
 - The different jobs recur at identical time points.




Major Cycle

- The major cycle of a set of tasks $ST=\{T1,T2,...,Tn\}$ is at least:
 - $LCM(p1,p2,...,pn)$
 - Holds even when tasks have arbitrary phasings.
 - Can be greater than LCM when F does not divide major cycle.

Minor Cycle (Frame)

- Each major cycle:
 - Usually contains an integral number of minor cycles (frames).
- Frame boundaries are marked:
 - Through interrupts generated from a periodic timer.

Major and Minor Cycle



- Cyclic scheduler runs in response to a tick event
- Red bar shows time to execute scheduler

A Typical Schedule

Frame	Task
F1	T2
F2	T3
F3	T2


Constructing a Schedule

- Construct static schedule for a **Major Cycle**
- Cyclic Executive repeats this schedule
- There may be resulting idle intervals
 - if so attempt to arrange so they occur periodically

Cyclic Schedule


f_1	tasks ₁
...	...
f_5	tasks ₅

Task = (r , e)	
T_1	= (4, 1)
T_2	= (5, 1.5)
T_3	= (20, 1)
T_4	= (20, 2)
H	= 20



Partitioning A Major Cycle into Frames


- Design steps:
 1. choose frame size,
 2. partition jobs into slices (if needed),
 3. place jobs/slices into frames.
 - At Frames boundaries :
 - **Cyclic executive** performs scheduling
 - There is no preemption within frame
- Major Cycle (Hyperperiod)



Cyclic Executive Example


- Job=(computation time, period, relative deadline)
- $A = (1, 10, 10)$
- $B = (3, 10, 10)$
- $C = (2, 20, 20)$
- $D = (8, 20, 20)$
 - Two slices $c(D1) = 2, c(D2) = 6$
- Major Cycle = 20msec
- We select frame size = 10 msec

Schedule Example



- Could we create a different schedule?
- Is it better to distribute the idle time?

Schedule Example




if Frame_no = 0 \rightarrow run A ; B ; C ; D1;

else \rightarrow run A ; B ; D2;

Frame_no = (Frame_no+1) mod 2;

Frame Size Constraints

1. Every job needs to start and complete within a frame
 - $f \geq \max(T_i), 1 \leq i \leq n$
2. Frame size f divides H (the Hyperperiod)
3. Between the release time and deadline of every job there is at least one frame



Minor Cycle (Frame)

- Each task is assigned to run in one or more frames.
- Frame size (F) is an important design parameter while using cyclic scheduler.
 - A selected frame size has to satisfy a few constraints.

Selecting an Appropriate Frame Size (F)

- Minimum scheduling overhead and chances of inconsistency:
 - F should be larger than each task size.
 - **Sets a lower bound.**
- Minimization of table size:
 - F should squarely divide major cycle.
 - **Allows only a few discrete frame size.**
- Satisfaction of task deadline:
 - Between the arrival of a task and its deadline:
 - At least one full frame must exist.
 - **Sets an upper bound**


Minimize Inconsistency

- Unless a job runs to completion:
 - Its partial results might be used by other jobs, leading to inconsistency
- To avoid scheduler overhead:
 - Selected frame size should be larger than execution time of each task.
 - Sets a lower bound for frame size.

Minimization of Table Size


- Unless the minor cycle squarely divides the major cycle:
 - Storing schedule for one major cycle would not be sufficient.
 - Schedules in the major cycle would not repeat:

- This would make the size of the table large.




Satisfaction of Task Deadline

- Between the arrival of a task and its deadline:
 - At least one full frame must exist.
- If there is not even a single frame:
 - The task would miss its deadline,
 - By the time it could be taken up for scheduling, the deadline could be imminent.



Satisfaction of Task Deadline


- The worst case for a task occurs when the task arrives just after a frame has started.



Satisfaction of Task Deadline

- The minimum separation of arrival time of t_i from a frame start:

- $GCD(F, p_i)$



- Thus, for all t_i

- $2F - GCD(F, p_i) \leq d_i$ must be satisfied

Frame Size Constraints

- (i) $f \geq \max(T_i)$ for all i
- (ii) $2f - \gcd(p_i, f) \leq D_i$
- (iii) $f^*(\text{LCM}(p_i)/f) = \text{LCM}(p_i)$

Selection of a Suitable Frame Size

- Several frame sizes may satisfy the constraints:
 - **Plausible frames.**
- A plausible frame size has been found:
 - Does not mean that the task set is schedulable.
- Also, the largest plausible frame size needs to be chosen:
 - **Scheduler overhead would be lower.**

Cyclic Scheduling Example 1

- Job=(computation time, period, relative deadline)
- $A = (1, 6, 6)$
- $B = (2, 8, 8)$
- Major cycle = 24
- Frame sizes: 2, 3, 4, 6, 12, 24
- For $F=2$
 - $2 \cdot F - \gcd(f, P_A) = 2 \cdot 2 - 2 \leq 6$ Acceptable
 - $2 \cdot F - \gcd(f, P_B) = 2 \cdot 2 - 2 \leq 8$ Acceptable

Example 1

- For $F=3$
 - $2 * F - \gcd(F, PA) = 6 - 3 \leq 6$ Acceptable
 - $2 * F - \gcd(F, PB) = 6 - 1 \leq 8$ Acceptable
- For $F=4$
 - $2 * F - \gcd(F, PA) = 8 - 2 \leq 6$ Acceptable
 - $2 * F - \gcd(F, PB) = 8 - 4 \leq 8$ Acceptable

But not enough frames available!

- We can choose $F=3$

$$\begin{aligned} A &= (1, 6, 6) \\ B &= (2, 8, 8) \end{aligned}$$

Example 2

- $A = (1, 10, 10)$
- $B = (3, 10, 10)$
- $C = (2, 20, 20)$
- $D = (8, 20, 20)$
- **Major Cycle**
 - $\text{Lcm}(10, 10, 20, 20) = 20$
- **Frame Size**
 - 1, 2, 4, 5, 10, 20 (must divide Major Cycle)
 - $f \geq \max\{1, 2, 3, 8\}$ (geq longest computation time)
 - f can be 10 or 20
- Not enough frames in a major cycle to run all the jobs